SOaC-Enterprise

A Scientific Framework for Autonomous Security Operations: The Convergence of Detection-as-Code, Explainable AI, and Programmable Response

Date: March 10, 2026

Version: 1.0 (Initial Release)

Classification: Public / Open-Source Research

Authors: SOaC Community / SOaC-Enterprise Research Group

Table of Contents

1. Executive Summary: The End of Reactive Triage

2. Introduction: The Entropy of Modern Security

3. The Four Pillars: A Distributed Intelligence Model

4. Scientific Foundations: Explainable AI (XAI) & Trust

5. Implementation: The Modular Threat Intelligence Package (MTIP)

6. Validation & Benchmarking: The Science of Trust

7. Case Study: The "Identity-to-Cloud" Pivot

8. Conclusion: The Future of Autonomous Defense

1. Executive Summary: The End of Reactive Triage

> *The modern Security Operations Center (SOC) is currently trapped in a "Linear Defense Paradox": while adversaries leverage automated, non-linear attack chains, defenders remain tethered to manual, linear triage workflows.*

SOaC-Enterprise (Security-Operations-as-Code) is a scientific framework designed to resolve this paradox. By decomposing security operations into four fundamental, programmable pillars—The Body (Telemetry), The Brain (Reasoning), The Purpose (Logic), and The Edge (Enforcement)—we introduce a distributed intelligence architecture that transforms the SOC from a reactive monitoring hub into an autonomous, code-driven defense engine.

This white paper details the mathematical and architectural foundations of the SOaC model, demonstrating how a vendor-agnostic, normalized telemetry layer combined with Explainable AI (XAI) can achieve sub-5-minute containment of sophisticated threat actors without the "black-box" risks associated with proprietary autonomous platforms.

2. Introduction: The Entropy of Modern Security

2.1 The Crisis of Complexity

In the last decade, the enterprise attack surface has expanded exponentially. The shift to hybrid-cloud architectures and the proliferation of SaaS applications have created "Identity Silos" that legacy SIEM and EDR platforms struggle to correlate. We identify three primary drivers of the current SOC crisis:

- **Alert Entropy:** The sheer volume of telemetry creates a signal-to-noise ratio that exceeds human cognitive limits.

- **The Skills Gap:** The requirement for analysts to be experts in KQL, SPL, Python, and multi-cloud IAM is unrealistic and unscalable.
- **Vendor Fragmentation:** Security logic is often "trapped" within proprietary vendor silos, preventing a unified, cross-stack response.

2.2 The SOaC Hypothesis

We hypothesize that security operations can be optimized by treating defense logic as First-Class Software. By applying the principles of DevOps—version control, CI/CD, and modularity—to detection and response, we can create a "Self-Healing Defense" that evolves at the speed of the adversary.

3. The Four Pillars: A Distributed Intelligence Model

3.1 The Body: Telemetry & The Physics of Ingestion

The "Body" represents the sensory nervous system of the enterprise. SOaC-Enterprise employs a Universal Normalization Layer based on the Elastic Common Schema (ECS). By mapping disparate vendor logs into a unified schema at the point of ingestion, we eliminate the "N+1 Integration Problem."

3.2 The Brain: Asynchronous Cognitive Reasoning

The "Brain" utilizes Large Language Models (LLMs) as Asynchronous Reasoning Engines. The cognitive architecture involves Contextual Injection, Chain-of-Thought (CoT) Processing, and a structured Verdict output (BLOCK/MONITOR).

3.3 The Purpose: Playbooks-as-Code (The Logic)

Playbooks are versioned YAML artifacts. This enables idempotency, peer review via Git Pull Requests, and portability across Azure, AWS, and On-Prem environments.

3.4 The Edge: Real-Time Enforcement Mechanics

The "Edge" translates the "Purpose" into API calls across Identity (Entra ID), Network (Firewalls/SGs), and Endpoint (EDR isolation) layers.

4. Scientific Foundations: Explainable AI (XAI) & Trust

Explainable AI (XAI): A primary barrier to autonomous security is the "Trust Gap." SOaC-Enterprise bridges this through Radical Transparency.

Every AI decision includes a ReasoningPath—a human-readable explanation of why the AI chose a specific action. This satisfies the transparency requirements of the EU AI Act and provides context for post-incident forensics.

5. Implementation: The Modular Threat Intelligence Package (MTIP)

ID	Package Name	Primary Focus
001	Identity Control Plane	Token Theft, PIM Abuse, MFA Bypass
002	Cloud Control Plane	AWS/Azure Root Abuse, IAM Shadowing
003	SaaS Pivot & Extortion	Salesforce, SAP, Google Workspace
006	GenAI & LLM Abuse	Prompt Injection, Data Exfiltration

6. Validation & Benchmarking: The Science of Trust

The Telemetry Replay Engine allows security teams to "Time Travel" by ingesting historical telemetry from known breaches to validate current configurations. Every package is benchmarked against the MITRE ATT&CK Framework with a detailed Coverage Heatmap.

7. Case Study: The "Identity-to-Cloud" Pivot

In a simulated Lazarus Group attack, Microsoft Sentinel detected anomalous token usage. The Brain identified a high-confidence identity pivot and issued a BLOCK verdict. Within 12 seconds, the Master Orchestrator revoked Entra ID sessions and restricted AWS IAM roles, containing the attack before exfiltration.

8. Conclusion: The Future of Autonomous Defense

SOaC-Enterprise is a Scientific Manifesto for the future of security operations. By treating defense as code and AI as an explainable reasoning engine, we empower enterprises to reclaim the advantage from the adversary.

© 2026 SOaC Community | SOaC-Enterprise Research Group